



ES215 - COA

Assignment 4

Siddhesh Kanawade, 20110199

[Github repository](#)

Question 1:

Q1) 5 stage pipeline with each stage taking one cycle.

We have to assume 5 different stages, hence we assume:

→ IF	ID	EX	MEM	WB
200ps	100ps	200ps	400ps	100ps

source: Assignment 3

→ Baseline is Single Cycle Interaction.

* Basic calculation:

$$\begin{aligned}T_{\text{net, single cycle}} &= \text{IF} + \text{ID} + \text{EX} + \text{EM} + \text{WB} \\&= 200 + 100 + 200 + 400 + 100 \\&= \underline{\underline{1000\text{ps}}}\end{aligned}$$

For N interactions →

$$T_{\text{net}} = 1000N$$

* Time taken for pipeline:

a) → 30% RAW & 20% branch dependency:

$$T_{\text{clock cycle, net}} = \max\{\text{Time by stages}\} = 400\text{ps}$$

For dependencies:

$$\begin{aligned}\text{CPI} &= \sum \text{CPI per inst.} \times \text{fraction} \\&= (1 - 0.2 - 0.3) \times 1 + 0.5 \times 4 + 0.2 \times 3 \\&= 0.5 \times 1 + 0.3 \times 4 + 0.2 \times 3 \\&= 0.5 + 1.2 + 0.6 \\&= \underline{\underline{2.3}}\end{aligned}$$

$$T_{exec} = IC \times CPI \times \text{Clock Cycle}$$

$$= N \times 2.3 \times 400$$

$$T_{exec} = \underline{\underline{920 N \text{ ps}}}$$

$$\therefore \text{Speedup} = \frac{T_{baseline}}{T_{pipeline}} = \frac{1000}{920} = \underline{\underline{1.087}}$$

b) \rightarrow 40% branch dependency:

$$CPI = (1 - 0.4) \times 1 + 0.4 (1 + \text{stall}_{branch})$$

$$= 0.6 \times 1 + 0.4 \times (1 + 2)$$

$$CPI = 0.6 + 1.2 = 1.8$$

$$\therefore T_{exec} = IC \times CPI \times \text{Clock Cycle Time}$$

$$= N \times 1.8 \times 400 \text{ ps}$$

$$= 720 N \text{ ps.}$$

$$\therefore \text{Speedup} = \frac{1000 N}{720 N} = \underline{\underline{1.3889}}$$

\Rightarrow Branch predictor with 80% accuracy.

a) 30% RAW, 20% branch

On calculation,

4% branch, 30% RAW, 66% Normal.

$$CPI = 0.66 \times 1 + 0.3 (1 + \text{stall}_{RAW}) + 0.04 (1 + \text{stall}_{branch})$$

$$= 0.66 \times 1 + 0.3 (1 + 3) + 0.04 (1 + 2)$$

$$= 0.66 + 1.2 + 0.12$$

$$= 1.98$$

$$T_{exec} = N \times 1.98 \times 400 = 792N \text{ ps.}$$

$$\therefore \text{Speedup} = \frac{1000}{792} = \underline{\underline{1.2626}}$$

⇒ Speedup increased from the previous value (without branch predictor).

⇒ 40% branch dependency:

⇒ 8% branch, 92% Normal.

$$\begin{aligned} \therefore CPI &= 0.92 \times 1 + 0.08 \times (1 + \text{stall}_{\text{branch}}) \\ &= 0.92 + 0.08(1+2) \\ &= 0.92 + 0.24 = \underline{\underline{1.16}} \end{aligned}$$

$$\therefore T_{exec} = N \times 1.16 \times 400 = 464N \text{ ps}$$

$$\therefore \text{Speedup} = \frac{1000}{464} = \underline{\underline{2.156}}$$



increased from previous value,
where there isn't branch predictor.

Question 2:

20% branch instruction.

Delayed branching with one delay slot.

Base CPI = 1.5 ... (No delay slots)

$$CPI_{base} = CPI_{nonbranch} \times f_{nonbranch} + CPI_{branch} \times f_{branch}$$

Thus,

$$1.5 = CPI_{nonbranch} \times 0.8 + CPI_{branch} \times 0.2 \quad \textcircled{1}$$

* 85% delayed slots are filled:

\therefore Fraction of branch instruction where delayed slot is filled $\Rightarrow 0.2 \times 0.85 = 0.17$.

\therefore fraction of branch instruction = $0.2 \times 0.15 = \underline{\underline{0.03}}$
(with delay slot = empty)

$$CPI_{new} = 0.8 \times CPI_{non-branch} + 0.2 \times CPI_{branch} - 0.17 \quad \textcircled{2}$$

from $\textcircled{1}$ & $\textcircled{2}$, we get

$$CPI_{new} = 1.5 - 0.17 = \underline{\underline{1.33}}$$

Question 3:

The code for the problem is present in the attached github repository. We could note that we have 3 nested for loops and we can interchange them maintaining the precision of the solution. Following are the combinations of the loop and their performance

Note: Time in Nanoseconds.

I, J, K

N	128	256	512
Iteration 1	11238176	87735141	738401854
Iteration 2	10315913	95228120	724739624
Iteration 3	10026547	88626850	716162352
Iteration 4	10444267	86697339	712669493
Iteration 5	10741363	114325474	734300589
Median	10444267	88626850	724739624

I, K, J

N	128	256	512
Iteration 1	9573563	59786247	420068011
Iteration 2	9467350	59430480	425174818
Iteration 3	11856721	58671657	419498810
Iteration 4	9306361	58621853	386511881
Iteration 5	9381577	59718417	420470088
Median	9467350	59430480	420068011

K, I, J

N	128	256	512
Iteration 1	9981769	58333633	422652882
Iteration 2	9305010	61117587	431034842

Iteration 3	8876656	59704946	426423398
Iteration 4	10132701	59349419	423859587
Iteration 5	8090361	62369342	422590950
Median	9305010	59704946	423859587

K, J, I

N	128	256	512
Iteration 1	13794999	102065793	814084227
Iteration 2	9774119	97619202	798926941
Iteration 3	13413033	103121638	852852350
Iteration 4	12716718	98763877	817270906
Iteration 5	10712025	108070460	811043671
Median	12716718	102065793	814084227

J, K, I

N	128	256	512
Iteration 1	12023385	103437377	812519930
Iteration 2	11741502	107421633	823129128
Iteration 3	13631846	101075464	828516363
Iteration 4	13316720	100282157	817735952
Iteration 5	12827583	103077558	797951845
Median	12827583	103077558	817735952

J, I, K

N	128	256	512
Iteration 1	9882507	89970305	750988468
Iteration 2	10625240	87557448	728962395

Iteration 3	10200303	90462679	743842431
Iteration 4	9880655	92903623	782598922
Iteration 5	12014748	87495620	724580132
Median	10200303	89970305	743842431


SUMMARY:

The summary of the above experiment is shown in the following table:


	128	256	512
IJK	10444267	88626850	724739624
IKJ	9467350	59430480	420068011
KIJ	9305010	59704946	423859587
KJI	12716718	102065793	814084227
JKI	12827583	103077558	817735952
JKI	10200303	89970305	743842431

Graphs(showing variation of N with permutation of i, j, k)


N = 128



N = 256



N = 512



SUMMARY

All the observations noted here have above graphs and experimental experience as basis

1. The overall **behavior** of graph for all values of N remain **similar**
2. **Minimum** time(best sequence) = **k, i, j**
3. **Maximum** time(worst sequence) = **j, k, i**
4. The performance of **k, i, j** and **i, k, j** are very close and in most cases the difference isn't significant. Same is true with **j, k, i** and **k, j, i**.
5. Most of the performance is determined by the inner most loop:
Performance(as inner most loop): $j > k > i$

The above graphs are self explanatory and above were the major conclusions.

Question 4:

The code for the problem is present in the attached github repository. We could note that we have 3 nested for loops and we can interchange them maintaining the precision of the solution. Following are the combinations of the loop and their performance

Note: Time in Seconds

I, J, K

N	128	256	512
Iteration 1	0.28224969199982297	2.073409698999967	16.996527972000877
Iteration 2	0.26032037300046795	2.0517706720002025	16.78452467199986
Iteration 3	0.26722570800029644	1.99435360499956	17.841820249000193
Median	0.26722570800029644	2.0517706720002025	16.996527972000877

I, K, J

N	128	256	512
Iteration 1	0.25056748799943307	1.8869987930002026	13.756732945000294
Iteration 2	0.2495645169992713	1.8789108379996833	14.168155080999895
Iteration 3	0.2748677259996839	1.9172490860000835	15.733258565000142
Median	0.25056748799943307	1.8869987930002026	14.168155080999895

K, I, J

N	128	256	512
Iteration 1	0.26016236700070294	1.9112885190006637	13.158623727999839
Iteration 2	0.2540488489994459	1.9509402309995494	14.424502025999573
Iteration 3	0.2537215259999357	1.9074726670005475	14.70270794499993
Median	0.2540488489994459	1.9112885190006637	14.424502025999573

K, J, I

N	128	256	512
Iteration 1	0.26719552600025054	1.995343726000101	17.234896037999533
Iteration 2	0.2630918090007981	2.048039387999779	17.18846589900022
Iteration 3	0.26402612299898465	2.1766023209993364	16.882021737000287
Median	0.26402612299898465	2.048039387999779	17.18846589900022

J, K, I

N	128	256	512
Iteration 1	0.28957370899843227	2.091274442998838	19.85446316099842
Iteration 2	0.2797036539996043	2.0989001909983926	19.695707141998355
Iteration 3	0.2809136839987332	2.059013183999923	19.81430137800089
Median	0.2809136839987332	2.091274442998838	19.81430137800089

J, I, K

N	128	256	512
Iteration 1	0.307991289999336	2.116191377999712	17.332238044000405
Iteration 2	0.27211802800047735	1.9963988079998671	15.595265014999313
Iteration 3	0.27281849100108957	2.0243031270001666	15.420025614999759
Median	0.27281849100108957	2.0243031270001666	15.595265014999313


SUMMARY:

The summary of the above experiment is shown in the following table:


	128	256	512
IJK	0.267225708	2.051770672	16.99652797
IKJ	0.250567488	1.886998793	14.16815508
KIJ	0.254048849	1.911288519	14.42450203
KJI	0.264026123	2.048039388	17.1884659
JKI	0.280913684	2.091274443	19.81430138
JIK	0.272818491	2.024303127	15.59526501

Graphs(showing variation of N with permutation of i, j, k)


N = 128



N = 256



N = 512



SUMMARY

All the observations noted here have above graphs and experimental experience as basis

1. The overall **behavior** of the graph for all values of N remains **similar**, but we observe significant deviation for N = 512.
2. The code execution in python takes a lot longer than code execution in C++, which was expected.
3. We observe a steep slope for N = 512, which is absent in rest two.
4. **Minimum** time(best sequence) = i, k, j
5. **Maximum** time(worst sequence) = j, k, i
6. The performance of k, i, j and i, k, j are very close and in most cases the difference isn't significant. Unlike for C++, the difference of performance between j, k, i and k, j, i, isn't negligible.
7. Most of the performance is determined by the inner most loop:

Performance(as inner most loop): $j > i > k$